



Denoising Diffusion Probabilistic Models

Giovani Tavares de Andrade

Instituto de Matemática e Estatística
(IME-USP)

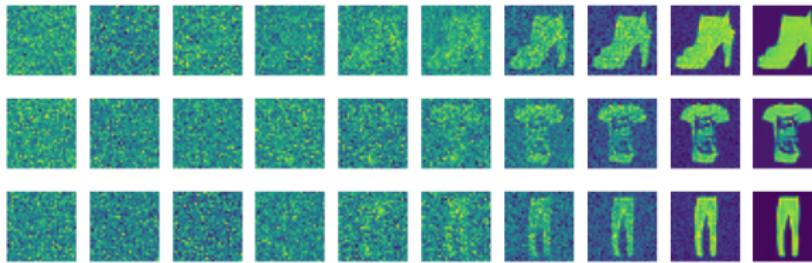
11 / 2025

Outline

- ① What are DDPMs
- ② Forward Process
- ③ Reverse Process
- ④ Training Algorithm
- ⑤ Sampling Algorithm
- ⑥ Bibliography

Overview

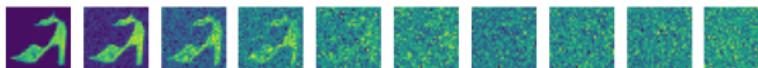
- Denoising Diffusion Probabilistic Models (DDPMs) generate images by iteratively removing noise from a signal until it looks like a real image.



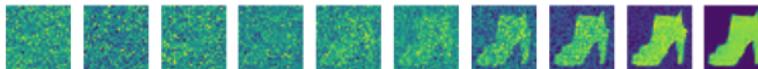
DDPMs Building Blocks

DDPMs are made of **two iterative processes** defined by different Markov Chains: a **forward** and a **reversion** process.

- ① **Forward Process:** gradually adds noise to a image by sampling from a normal distribution according to a Markov Chain



- ② **Reverse Process:** removes noise by sampling from a normal distribution according to another Markov Chain



DDPMs Building Blocks

① Forward Process' Distribution

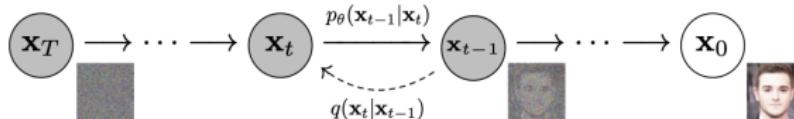
$$q(x_t|x_{t-1}) := \mathcal{N}(x_t; \sqrt{1 - \beta_t} \times x_{t-1}; \beta_t I) \quad (1)$$

② Reverse Process' Distribution

$$p(x_T) = \mathcal{N}(x_T; 0; 1) \quad (2)$$

$$p_\theta(x_{0:T}) := p(x_T) \prod_{t=1}^T p_\theta(x_{t-1}|x_t) \quad (3)$$

$$p_\theta(x_{t-1}|x_t) := \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t); \Sigma_\theta(x_t, t)) \quad (4)$$



Forward Process Distribution

- The forward distribution (\mathbf{q}) is a Normal Distribution with mean as a function of \mathbf{x}_0 :

$$\mathbf{q}(\mathbf{x}_t | \mathbf{x}_{t-1}) := \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \times \mathbf{x}_{t-1}; \beta_t I) \quad (5)$$

(6)

$$\boxed{\mathbf{x}_t} \sim \mathcal{N}\left(\sqrt{\bar{\alpha}_t} \boxed{\mathbf{x}_0}, (\mathbf{1} - \bar{\alpha}_t)\mathbf{I}\right)$$

- $\bar{\alpha}_t$ is called the cumulative noise and is proportional to the product of the transitions distributions that gradually turns \mathbf{x}_0 into \mathbf{x}_t .
- The forward distribution let's us produce a noisy signal \mathbf{x}_t from the original \mathbf{x}_0 in a single step.

Forward Process in PyTorch

```
1  class DDPMsScheduler(nn.Module):
2      def __init__(self, num_time_steps: int=1000):
3
4          super().__init__()
5          self.beta = torch.linspace(1e-4, 0.02, num_time_steps, requires_grad=False)
6
7          # Alpha is defined as 1 - beta, i.e., 1 - Variance at step t
8          alpha = 1 - self.beta
9
10         # The variances are held constant during training, requires_grad = False
11         self.alpha = torch.cumprod(alpha, dim=0).requires_grad_(False)
12
13     def forward(self, t):
14         return self.beta[t], self.alpha[t]
15
```

Forward Process Sampling in PyTorch

Given an input tensor \mathbf{x}_0 one can sample from $\mathbf{q}(\mathbf{x}_t | \mathbf{x}_0)$ by making the attribution:

```
1 # Random sample from dataset
2 idx = random.randint(0, len(dataset) - 1)
3 x0, y0 = dataset[idx]
4
5
6 # Sampling 10 signal with increasing noise levels
7 for t in range(1, 10):
8     scheduler = DDPMscheduler(num_time_steps=t + 1)
9
10    # Cumulative noise up until step t
11    beta, alpha = scheduler(t)
12
13    e = torch.randn_like(x0, requires_grad=False)
14    x_t = (torch.sqrt(alpha)*x0) + (torch.sqrt(1-alpha)*e)
15
16
17
```

Reverse Process' Transitions Distribution

- The reverse distribution is a Normal Distribution with mean and variance as functions of the noisy signal \mathbf{x}_t and the amount of noise that has been added to it t .
-

$$p_\theta(x_{t-1}|x_t) := \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t); \Sigma_\theta(x_t, t)) \quad (7)$$

- What we want is to have a model that is capable of sampling from \mathbf{p} so that \mathbf{x}_{t-1} is indeed a less noisy version of \mathbf{x}_t , i.e., a model that can reverse what was done in the forward process.

Reverse Process' Prior

- Ideally, the Reverse Process would enable us to sample from \mathbf{x}_0 directly:

$$p(\mathbf{x}_0) = \int p(\mathbf{x}_0, \mathbf{x}_{1:T}) d\mathbf{x}_{1:T} \quad (8)$$

(9)

- Unfortunately, $p(\mathbf{x}_0)$ is very complex due to its multidimensionality. It is **intractable**.
- That is why the training of DDPMs never optimizes $p(\mathbf{x}_0)$ directly. Instead, what is optimized is its lower bound.

Reverse Process' Transitions Distribution

-

$$p(x_T) = \mathcal{N}(x_T; 0; 1) \quad (10)$$

$$p_{\theta}(x_{0:T}) := p(x_T) \prod_{t=1}^T p_{\theta}(x_{t-1}|x_t) \quad (11)$$

$$p_{\theta}(x_{t-1}|x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t); \Sigma_{\theta}(x_t, t)) \quad (12)$$

- What we ultimately want is to have a \mathbf{x}_0 that is as likely as possible. We can marginalize \mathbf{x}_0 using the latent variables $\mathbf{x}_{1:T}$, i.e., by using the noisy images.

Evidence Lower Bound (ELBO)

- DDPMs are not trained to sample from $p(\mathbf{x}_0)$, but instead to maximize its lower bound ELBO (L)
-

$$\log[p_\theta(\mathbf{x}_0)] = \log \int_{\mathbf{x}_{1:T}} p(\mathbf{x}_0, \mathbf{x}_{1:T}) d\mathbf{x}_{1:T} \quad (13)$$

$$\log[p_\theta(\mathbf{x}_0)] = \log \int_{\mathbf{x}_{1:T}} p(\mathbf{x}_0, \mathbf{x}_{1:T}) \frac{q(\mathbf{x}_{1:T}|\mathbf{x}_0)}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} d\mathbf{x}_{1:T} \quad (14)$$

by definition, (15)

$$\mathbb{E}_q[f(\mathbf{x}_{1:T})] = \int q(\mathbf{x}_{1:T}|\mathbf{x}_0) f(\mathbf{x}_{1:T}) d\mathbf{x}_{1:T} \quad (16)$$

$$\implies \log[p_\theta(\mathbf{x}_0)] = \log(\mathbb{E}_q \left[\frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right]) \quad (17)$$

the Jensen's inequality tells us (18)

$$f(\mathbb{E}(\mathbf{X})) \geq \mathbb{E}(f(\mathbf{X})) \quad (19)$$

for any concave function f . \log is concave, hence: (20)

$$\log[p_\theta(\mathbf{x}_0)] \geq \mathbb{E}_q \left[\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \quad (21)$$

We define the Evidence Lower Bound L as: (22)

$$L := \mathbb{E}_q \left[-\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \quad (23)$$

Evidence Lower Bound (ELBO)

- DDPMs are not trained to sample from $p(\mathbf{x}_0)$, but instead to maximize its lower bound ELBO (L)
- With more algebraic manipulation and using the fact that both the forward and reverse processes are Markov Chains, one can derive the following equation:

$$L := \mathbb{E}_q \left[-\log \frac{p(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T}|\mathbf{x}_0)} \right] \quad (25)$$

$$L = \mathbb{E}_q \left[-\log \frac{p(\mathbf{x}_T)}{q(\mathbf{x}_T|\mathbf{x}_0)} \right] - \log p_\theta(\mathbf{x}_0|\mathbf{x}_1) - \mathbb{E}_q \left[\sum_{t=2}^T D_{KL} [q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) || p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)] \right] \quad (26)$$

- **We can conclude that maximizing ELBO (L) is equivalent to minimizing the KL-Divergence between $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ and $p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)$, i.e., minimizing the divergence between the forward and reverse processes' distributions.**

Defining $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \frac{q(\mathbf{x}_t | \mathbf{x}_{t-1}, \mathbf{x}_0) q(\mathbf{x}_{t-1} | \mathbf{x}_0)}{q(\mathbf{x}_t | \mathbf{x}_0)} \quad (27)$$

we know the q distribution from Definition, hence (28)

$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$ is a product of known Gaussians over another known Gaussian (29)

$$\mu_q(\mathbf{x}_t, \mathbf{x}_0) = \frac{(1 - \bar{\alpha}_{t-1})\sqrt{\alpha_t}\mathbf{x}_t + (1 - \alpha_t)\sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_0}{(1 - \bar{\alpha}_t)} \quad (30)$$

$$\Sigma_q(t) = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{I} \quad (31)$$

$$\implies q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_q(\mathbf{x}_t, \mathbf{x}_0); \Sigma_q(t)) \quad (32)$$

Defining $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t); \Sigma_\theta(t)) \quad (33)$$

from DDPM's paper: (34)

$$\Sigma_\theta(t) = \Sigma_q(t) \quad (35)$$

we are only left with the distribution's mean μ_θ (36)

$$\implies p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t); \Sigma_q(t)) \quad (37)$$

Computing the KL Divergence Between q and p_θ

- Now we know we are trying to compute the KL Divergence between two Gaussians with the exact same variance.
- For that, there is the following result that arrives from the definition of such divergence

$$d_1(x) = \mathcal{N}(\mu_1, \sigma^2) \quad (38)$$

$$d_2(x) = \mathcal{N}(\mu_2, \sigma^2) \quad (39)$$

The KL divergence $D_{KL}(d_1 \| d_2)$ is given by: (40)

$$D_{KL}(d_1 \| d_2) = \frac{(\mu_1 - \mu_2)^2}{2\sigma^2} \quad (41)$$

Hence, (42)

$$\mathbf{D}_{KL}[q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) || p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)] = \mathbf{D}_{KL}(\mathcal{N}(\mathbf{x}_{t-1}; \mu_q(\mathbf{x}_t, \mathbf{x}_0); \Sigma_q(t)), \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t); \Sigma_q(t))) \quad (43)$$

$$= \frac{1 - \bar{\alpha}_t}{2(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})} \| (\mu_q - \mu_\theta)_2^2 \| \quad (44)$$

- We just need to minimize the difference between the means of the reverse and forward processes' distributions, i.e., minimize $\| (\mu_q - \mu_\theta)_2^2 \|$.**

Defining The Model's Prediction

- We can use the prediction of our model as the forward process' mean
-

$$\mu_q(\mathbf{x}_t, \mathbf{x}_0) = \frac{(1 - \bar{\alpha}_{t-1})\sqrt{\alpha_t}\mathbf{x}_t + (1 - \alpha_t)\sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_0}{(1 - \bar{\alpha}_t)} \quad (45)$$

we can define the prediction (46)

$$\mu_\theta(\mathbf{x}_t) := \hat{\mu}_q(\mathbf{x}_t, \mathbf{x}_0) \quad (47)$$

$$= \frac{(1 - \bar{\alpha}_{t-1})\sqrt{\alpha_t}\mathbf{x}_t + (1 - \alpha_t)\sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_\theta}{(1 - \bar{\alpha}_t)} \quad (48)$$

$$\implies \mathbf{D}_{\mathbf{KL}}(\mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t); \Sigma_q(t)), \mathcal{N}(\mathbf{x}_{t-1}; \mu_q(\mathbf{x}_t, \mathbf{x}_0); \Sigma_q(t))) \quad (49)$$

$$= \frac{(1 - \bar{\alpha}_t)(\bar{\alpha}_{t-1})}{2(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})} \|(\mathbf{x}_\theta - \mathbf{x}_0)_2^2\| \quad (50)$$

- **Theoretically, this equation could be used as the loss function directly, but we can rewrite the images \mathbf{x}_θ and \mathbf{x}_0 in function of the Gaussian noises.**

Defining The Model's Prediction

- We can rewrite \mathbf{x}_t and \mathbf{x}_0 as functions of the added gaussian noises
-

$$q(\mathbf{x}_t | \mathbf{x}_0) := \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} \times x_0; (1 - \bar{\alpha}_t)\mathbb{I}) \quad (51)$$

which let's us write

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon \quad (53)$$

$$\Rightarrow \mathbf{x}_0 = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon}{\sqrt{\bar{\alpha}_t}} \quad (54)$$

for a Standard Gaussian Noise ϵ . (55)

We can now define our prediction $\hat{\epsilon} = \epsilon_\theta$ (56)

$$\mathbf{x}_\theta = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_\theta}{\sqrt{\bar{\alpha}_t}} \quad (57)$$

$$\Rightarrow \frac{(1 - \bar{\alpha}_t)(\bar{\alpha}_{t-1})}{2(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})} \|(\mathbf{x}_\theta - \mathbf{x}_0)_2^2\| = \frac{(1 - \bar{\alpha}_t)(\bar{\alpha}_{t-1})}{2(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})} \frac{(1 - \alpha_t)^2}{(1 - \bar{\alpha}_t)\alpha_t} \|(\epsilon_\theta - \epsilon)_2^2\| \quad (58)$$

- **The DDPM paper authors mention that optimizing $\|(\epsilon_\theta - \epsilon)_2^2\|$ without the scaling factor with the cumulative noise α_t is enough.**

Noise Predictor Training

- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ ▷ Sample image from training set
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$ ▷ Sample the step of the Forward Process Markov Chain
- 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ ▷ Sample standard gaussian noise to be added to the input
- 5: $\mathbf{x}_t = \sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon$ ▷ Forward Process/ Generating Noisy Image
- 6: Take Gradient Descent Step on $\nabla_{\theta}(\|\epsilon - \epsilon_{\theta}(\mathbf{x}_t, t)\|)$
- 7: **until** converged

Sampling Algorithm Derivation

-

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_\theta(\mathbf{x}_t); \Sigma_\theta(t)) \quad (59)$$

$$\mu_\theta(\mathbf{x}_t) = \frac{(1 - \bar{\alpha}_{t-1})\sqrt{\alpha_t}\mathbf{x}_t + (1 - \alpha_t)\sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_\theta}{(1 - \bar{\alpha}_t)} \quad (60)$$

$$\mathbf{x}_\theta = \frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t}\epsilon_\theta}{\sqrt{\bar{\alpha}_t}} \quad (61)$$

$$\implies \mu_\theta(\mathbf{x}_t) = \frac{\mathbf{x}_t}{\sqrt{\alpha_t}} - \frac{(1 - \alpha_t)(\sqrt{1 - \bar{\alpha}_t})}{(1 - \bar{\alpha}_t)(\sqrt{\alpha_t})}\epsilon_\theta = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}}\epsilon_\theta \right) \quad (62)$$

$$\Sigma_\theta(t) = \Sigma_q(t) = \frac{(1 - \alpha_t)(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \mathbf{I} \quad (63)$$

We now have defined \mathbf{x}_{t-1} 's mean and variance given \mathbf{x}_t which let's us write it as
(64)

$$\mathbf{x}_{t-1} = \mu_\theta(\mathbf{x}_t) + \sqrt{\Sigma_\theta(t)}\mathbf{z} \quad (65)$$

$$\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \quad (66)$$

Sampling Algorithm

```
1: repeat
2:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$                                 ▷ Sample random noisy image
3:  $T \sim \text{Uniform}(\{1, \dots, 1000\})$                 ▷ Sample random length of the Denoising Chain
4: for  $t = T, \dots, 1$  do
5:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$  else  $\mathbf{z} = \mathbf{0}$ 
6:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sqrt{\Sigma_\theta(t)} \mathbf{z}$     ▷ Sampling  $\mathbf{x}_{t-1}$  from  $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ 
7: end for
8: return  $\mathbf{x}_0$ 
```

References I

- [1] Ho, J., Jain, A., and Abbeel, P. (2020). *Denoising Diffusion Probabilistic Models*. <https://arxiv.org/pdf/2006.11239>
- [2] Sohl-Dickstein, J., Weiss, E., Maheswaranathan, N., and Ganguli, S. (2015). *Deep Unsupervised Learning using Nonequilibrium Thermodynamics*. <https://arxiv.org/abs/1503.03585>
- [3] 3Blue1Brown. (n.d.). *Neural Networks [7.8]: Deep Learning - Variational Bound (YouTube)*.
<https://www.youtube.com/watch?v=pStDscJh2Wo>
- [4] Jake Tae. (2021). *A Step Up with Variational Autoencoders*.
<https://jaketae.github.io/study/vae/>
- [5] Jake Tae. (2021). *From ELBO to DDPM*.
<https://jaketae.github.io/study/elbo/>

References II

- [6] Yang, X. (2017). *Understanding the Variational Lower Bound*. <https://xyang35.github.io/2017/04/14/variational-lower-bound/>
- [7] AI Coffee Break with Letitia. (n.d.). *Denoising Diffusion Probabilistic Models — DDPM Explained (YouTube)*.
<https://www.youtube.com/watch?v=H45lF4sUgiE&t=880s>

Hora da Implementação :)